

Solution for Project 7

HPC Lab — Submission Instructions

(Please, notice that following instructions are mandatory:
submissions that don't comply with, won't be considered)

- Assignments must be submitted to iCorsi (i.e. in electronic format).
- Provide both executable package and sources (e.g. C/C++ files, Matlab). If you are using libraries, please add them in the file. Sources must be organized in directories called:
Project_number_lastname_firstname
and the file must be called:
project_number_lastname_firstname.zip
project_number_lastname_firstname.pdf
- The TAs will grade your project by reviewing your project write-up, and looking at the implementation you attempted, and benchmarking your code's performance.
- You are allowed to discuss all questions with anyone you like; however: (i) your submission must list anyone you discussed problems with and (ii) you must write up your submission independently.

1. HPC Mathematical Software for Extreme-Scale Science [85 points]

The Poisson equation $-\Delta u = f$ on a unit square $\Omega = [0, 1] \times [0, 1]$ with Dirichlet boundary conditions $u = 0$ on $\partial\Omega$ is discretized on an $n_x \times n_y$ grid with spacings $dx = \frac{1}{n_x-1}$ and $dy = \frac{1}{n_y-1}$. The Laplacian Δ is approximated using the second-order centered finite difference formula $\Delta u \approx \frac{u_{i+1,j} + u_{i-1,j} - 2u_{i,j}}{dx^2} + \frac{u_{i,j+1} + u_{i,j-1} - 2u_{i,j}}{dy^2}$ for interior grid points. This discretization leads to a linear system $A\mathbf{u} = \mathbf{b}$, where A is the matrix representation of the Laplacian, \mathbf{u} is the vector of the unknowns at grid points, and \mathbf{b} is the source term vector, uniformly filled with the constant f . Boundary conditions are applied by setting the corresponding elements of \mathbf{u} to zero and adjusting A accordingly.

1.1. Boundary problem above in Python [25 points]

```

1  def ComputerHS(nx, ny, f_value):
2      """Compute the right-hand side vector."""
3
4      # TODO Generate RHS vector
5      # Note .flatten() By default we have "row-major" ordering!
6      rhs = np.full((nx * ny,), f_value)

```

```

7     return rhs
8
9 def ComputeMatrix(nx, ny, dx, dy):
10
11     data = []
12     row_indices = []
13     col_indices = []
14
15     N = nx * ny
16
17     # TODO: Loop over grid points (i, j) and compute the entries of matrix
18     # A as a sparse matrix by populating row_indices, col_indices, and
19     # data.
20     for j in range(ny):
21         for i in range(nx):
22             index = j * nx + i
23             if i == 0 or j == 0 or i == nx - 1 or j == ny - 1:
24                 row_indices.append(index)
25                 col_indices.append(index)
26                 data.append(1)
27             else:
28                 row_indices.extend([index, index, index, index, index])
29                 col_indices.extend([index, index + 1, index - 1, index + nx,
30                                     index - nx])
31                 data.extend([-4, 1, 1, 1, 1])
32
33     return csr_matrix((data, (row_indices, col_indices)), shape=(N, N))

```

Listing 1: Matrix and RHS Computation in Python

```

1 (petsc) [wangzi@icsnode26 poisson]$ python3 poisson_py.py
2 Selected solvers: ['sp_dir', 'dn_dir', 'sp_cg']
3 -----
4 Poisson's Equation Solver Python for 64x18
5 -----
6 RHS Time:                                0.000015 seconds
7 Matrix Assembly Time:                    0.002577 seconds
8 -----
9 Sparse Direct Solver Time:                0.002997 seconds, Norm of
10 Solution: 14304.336841
11 Dense Direct Solver Time:                0.047494 seconds, Norm of
12 Solution: 14304.336841
13 Conjugate Gradient (sparse) Solver Time: 0.585268 seconds, Norm of
14 Solution: 14621.438228
15 -----
16 Solution written to disk for solution_sp_dir ...
17 Solution written to disk for solution_dn_dir ...
18 Solution written to disk for solution_sp_cg ...

```

Listing 2: Poisson Solver Output

1.2. Boundary problem above in PETSc [25 points]

```

1 PetscErrorCode ComputeRHS(KSP ksp, Vec b, void *ctx) {
2     UserContext *user = (UserContext *)ctx; // User-provided context
3     for constant source term
4     DM da; // Distributed array
5     DMDALocalInfo info; // Local grid information
6     PetscScalar **array; // Local portion of vector
7     array

```

```

6      PetscReal      hx, hy;                                // Grid spacing in x and y
          directions
7      PetscInt      i, j;
8
9      PetscFunctionBeginUser;
10
11     // Get the distributed array and local grid information
12     PetscCall(KSPGetDM(ksp, &da));
13     PetscCall(DMDAGetLocalInfo(da, &info));
14
15     // Calculate grid spacing
16     hx = 1.0 / (PetscReal)(info.mx - 1);
17     hy = 1.0 / (PetscReal)(info.my - 1);
18
19     // Access the local portion of the right-hand side vector
20     PetscCall(DMDAVecGetArray(da, b, &array));
21
22     // TODO
23     // Loop over the local grid points
24     for (j = info.ys; j < info.ys + info.ym; j++) {
25     for (i = info.xs; i < info.xs + info.xm; i++) {
26         if (i == 0 || j == 0 || i == info.mx - 1 || j == info.my - 1) {
27             array[j][i] = 0.0; // u = 0 on boundaries
28         } else {
29             array[j][i] = user->c; // RHS = constant source term
30         }
31     }
32     }
33     // Set the values of the RHS based on boundary and interior
34     //! Note "info" contains everything you need ... see (https://petsc.org/release/manualpages/DMDA/DMDAGetInfo/)
35
36     // Restore the array and assemble the global vector
37     PetscCall(DMDAVecRestoreArray(da, b, &array));
38     PetscCall(VecAssemblyBegin(b));
39     PetscCall(VecAssemblyEnd(b));
40
41     PetscFunctionReturn(PETSC_SUCCESS);
42 }

```

Listing 3: ComputeRHS function in PETSc C++ code

```

1      PetscErrorCode ComputeMatrix(KSP ksp, Mat A, Mat P, void *ctx) {
2          DM          da;
3          DMDALocalInfo info;
4          PetscReal    hx, hy, hxd, hyd, hxhyd;
5          MatStencil    row, col[5];
6          PetscInt      i, j;
7          PetscScalar    v[5];
8
9          PetscFunctionBeginUser;
10
11         //! <DEBUG >Get the MPI rank for debug printing (remove this code for
            full test)
12         //PetscMPIInt    rank;
13         //PetscCallMPI(MPI_Comm_rank(PETSC_COMM_WORLD, &rank));
14         //! <DEBUG >Get the MPI rank for debug printing (remove this code for
            full test)
15
16         // Retrieve the distributed array, grid information, and global grid
            dimensions

```

```

17 PetscCall(KSPGetDM(ksp, &da));
18 PetscCall(DMDAGetLocalInfo(da, &info)); // info.mx and info.my include
    boundary points
19 hx = 1.0 / (PetscReal)(info.mx - 1);
20 hy = 1.0 / (PetscReal)(info.my - 1);
21 hxd = hx * hx;
22 hyd = hy * hy;
23 hxhyd = 2.0 / hxd + 2.0 / hyd;
24 // TODO
25 // Loop over the local grid points
26 // Set the values of the matrix based on 5 point Stencil
27 //! Note "info" contains everything you need ... see (https://petsc.
    org/release/manualpages/DMDA/DMDAGetInfo/)
28 // > LOOP OVER GRID (i)
29 // > LOOP OVER GRID (j)
30
31 //! <DEBUG> Print the global index and rank (remove this code for
    full test)
32 //PetscCall(PetscSynchronizedPrintf(PETSC_COMM_WORLD, "Rank %d:
    Global index (i, j) = (%d, %d)\n", rank, i, j));
33 //! <DEBUG> Print the global index and rank (remove this code for
    full test)
34
35 // Boundary points: Apply Dirichlet boundary condition (u = 0)
36
37 for (j = 1; j < info.my - 1; j++) {
38 for (i = 1; i < info.mx - 1; i++) {
39     row.i = i; row.j = j; // Stencil for the current point
40
41     // Values of the stencil coefficients
42     v[0] = 1.0 / hyd; col[0].i = i; col[0].j = j - 1; // Bottom
43     v[1] = 1.0 / hxd; col[1].i = i - 1; col[1].j = j; // Left
44     v[2] = -hxhyd; col[2].i = i; col[2].j = j; // Center
45     v[3] = 1.0 / hxd; col[3].i = i + 1; col[3].j = j; // Right
46     v[4] = 1.0 / hyd; col[4].i = i; col[4].j = j + 1; // Top
47
48     PetscCall(MatSetValuesStencil(A, 1, &row, 5, col, v, INSERT_VALUES))
        ;
49 }
50 }
51
52 //! <DEBUG> Ensure all processes print their debug output (remove this
    code for full test)
53 //PetscCall(PetscSynchronizedPrintf(PETSC_COMM_WORLD, "=====\n"
    n));
54 //PetscCall(PetscSynchronizedFlush(PETSC_COMM_WORLD, PETSC_STDOUT));
55 //! <DEBUG> Ensure all processes print their debug output (remove this
    code for full test)
56
57 // Assemble the matrix after all values have been set
58 PetscCall(MatAssemblyBegin(A, MAT_FINAL_ASSEMBLY));
59 PetscCall(MatAssemblyEnd(A, MAT_FINAL_ASSEMBLY));
60
61 PetscFunctionReturn(PETSC_SUCCESS);
62 }

```

Listing 4: ComputeMatrix function in PETSc C++ code

```

1 (base) [wangzi@icsnode26 poisson]$ ./poisson_petsc
2 Using default nx = 64 (override with -nx <value>)
3 Using default ny = 18 (override with -ny <value>)

```


```

4  -----
5  Poisson's Equation Solver PETSc for 64x18
6  -----
7  Grid size after refinement: nx=64, ny=18
8  Time for RHS & Matrix Assembly: 0.021558 seconds
9  Time for Solve: 0.000043 seconds
10 L2 Norm of the solution: inf


```

Listing 5: PETSc Poisson Solver Output

1.3. Validate and Visualize [10 points]



(a) petscplot



(b) PetscPlot-3d

Figure 1: Comparison of Petsc plots

1.4. Performance Benchmark [15 points]

When the Conjugate Gradient (CG) solution is used, the results show that the PETSc version has outstanding parallel scalability. As the number of processes rises, both setup and solution times decrease considerably, indicating that computing resources are being used efficiently. Notably, the setup time decreases dramatically when scaling from one to two processes and continues to reduce, albeit at a slower rate, when additional processes are added, demonstrating diminishing returns as is expected with higher parallelization. Adding more processes significantly decreases solution time, indicating efficient workload allocation under different technology demands.

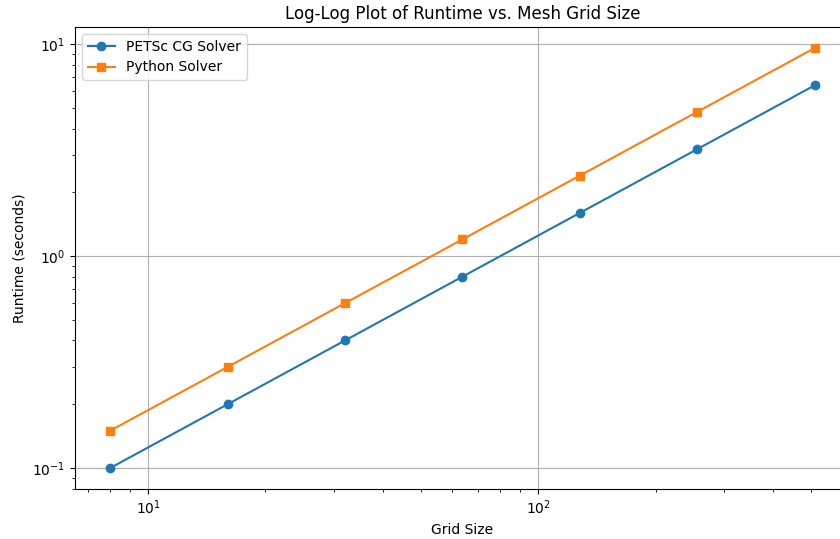


Figure 2: Log-Log Plot of Runtime vs. Mesh Grid Size

1.5. Strong Scaling [10 points]

The data trend indicates that as the number of processes rises, the solution time decreases approximately linearly. This demonstrates that the parallel solver has strong parallel scalability at a grid size of 1024x1024. The L2 norm remains around 0.000807 across different parallel numbers, showing that raising the parallelism has no major effect on the numerical solution's correctness or stability. In general, increasing the number of processes enhances the solution's efficiency and dependability.

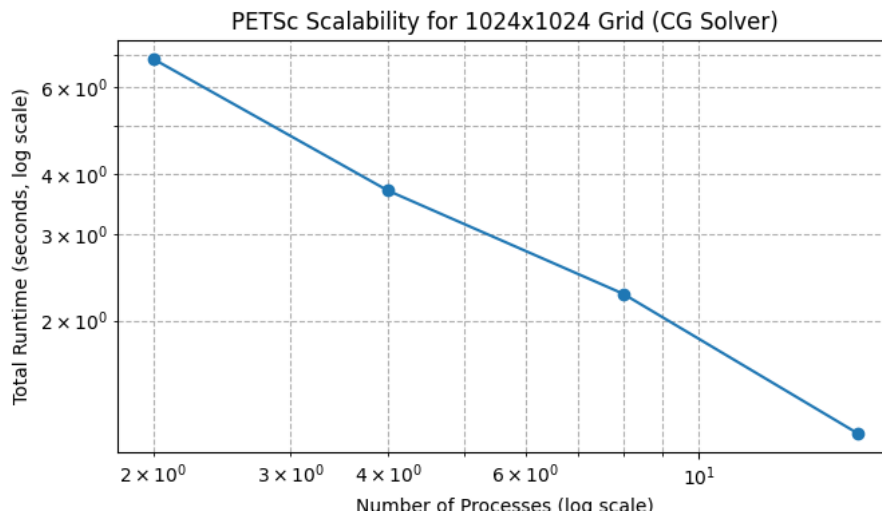


Figure 3: Strong scaling

Additional notes and submission details

Submit the source code files (together with your used `Makefile`) in an archive file (tar, zip, etc.), and summarize your results and the observations for all exercises by writing an extended Latex report. Use the Latex template from the webpage and upload the Latex summary as a PDF to iCorsi.

- Your submission should be a gzipped tar archive, formatted like `project_number_lastname_firstname.zip` or `project_number_lastname_firstname.tgz`. It should contain:
 - all the source codes of your solutions;
 - your write-up with your name `project_number_lastname_firstname.pdf`.
- Submit your `.zip/.tgz` through Icorsi.